Forecast, publication and outcome pipelines

Detailed control flow: freeze evidence before results; distinguish internal capture from observed public availability.

Forecast and publication

1. Fetch as-of sources
2. Generate base -> serve once
3. Journal before dedupe
4. Gate exact generation
5. Publish + verify hashes
6. Export + durable ACK

Failure path before publication

- Unknown status -> quarantine its identifiable row.
- Stale source or hash mismatch -> reject with reason.
- Keep R102 checks and R106 archive protection.

Failure path after publication

- Engine down -> preserve outbox; retry same IDs.
- Conflicting payload -> reject; never overwrite.
- ACK only after commit; alert before kickoff.

Outcome and learning

1. Authoritative FINAL
2. Append revision chain
3. Pin exact cohort/revisions
4. Chronological evaluation
5. Shadow candidate
6. NFL gate + owner review

Three clocks: source availability -> forecast creation -> internal/database observation. A later import cannot prove earlier capture. Public availability needs a separate deployment/content-hash observation.

Sampling boundary

Pool rungs, MY offers and Game/Week cards are distinct populations. Compare the same declared cohort; copies are not independent samples.

Scheduling and recovery

Actions completion + manual dispatch; cron is best effort. Durable leases/checkpoints recover repeated or killed jobs.