One learning engine, separate serving authority

Purpose and boundaries: preserve what was predicted, learn from trustworthy outcomes, and improve only with evidence.

Verified current foundation | Proposed integration | Incomplete or requires approval

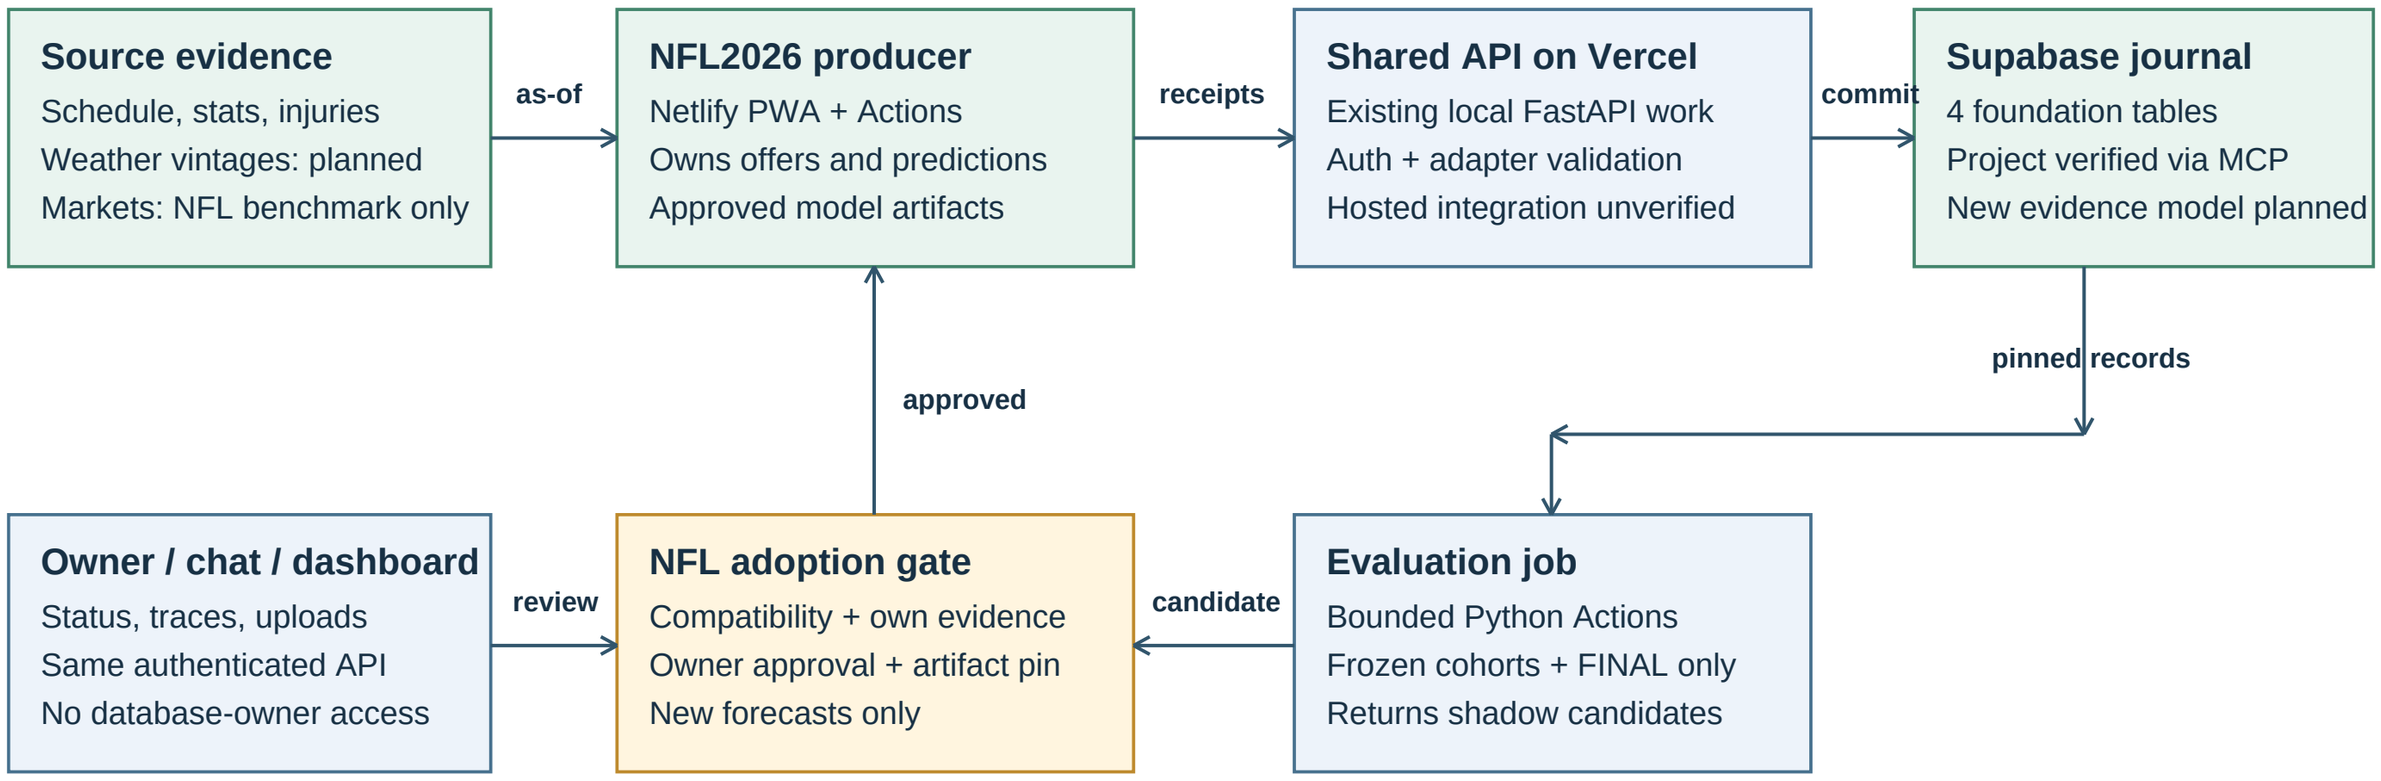

Current truth

Operator token is configured per owner. Supabase management access and foundation tables are verified. Neither proves a working NFL feedback loop.

Non-goals

No bets or trades, no automatic promotion, no guaranteed parlay wins, no new frontend framework, and no reconstruction of missing historical evidence.

Components, contracts and persistence

Mid-level design: the database owns durable observations; NFL owns the contract users saw and what serves next.

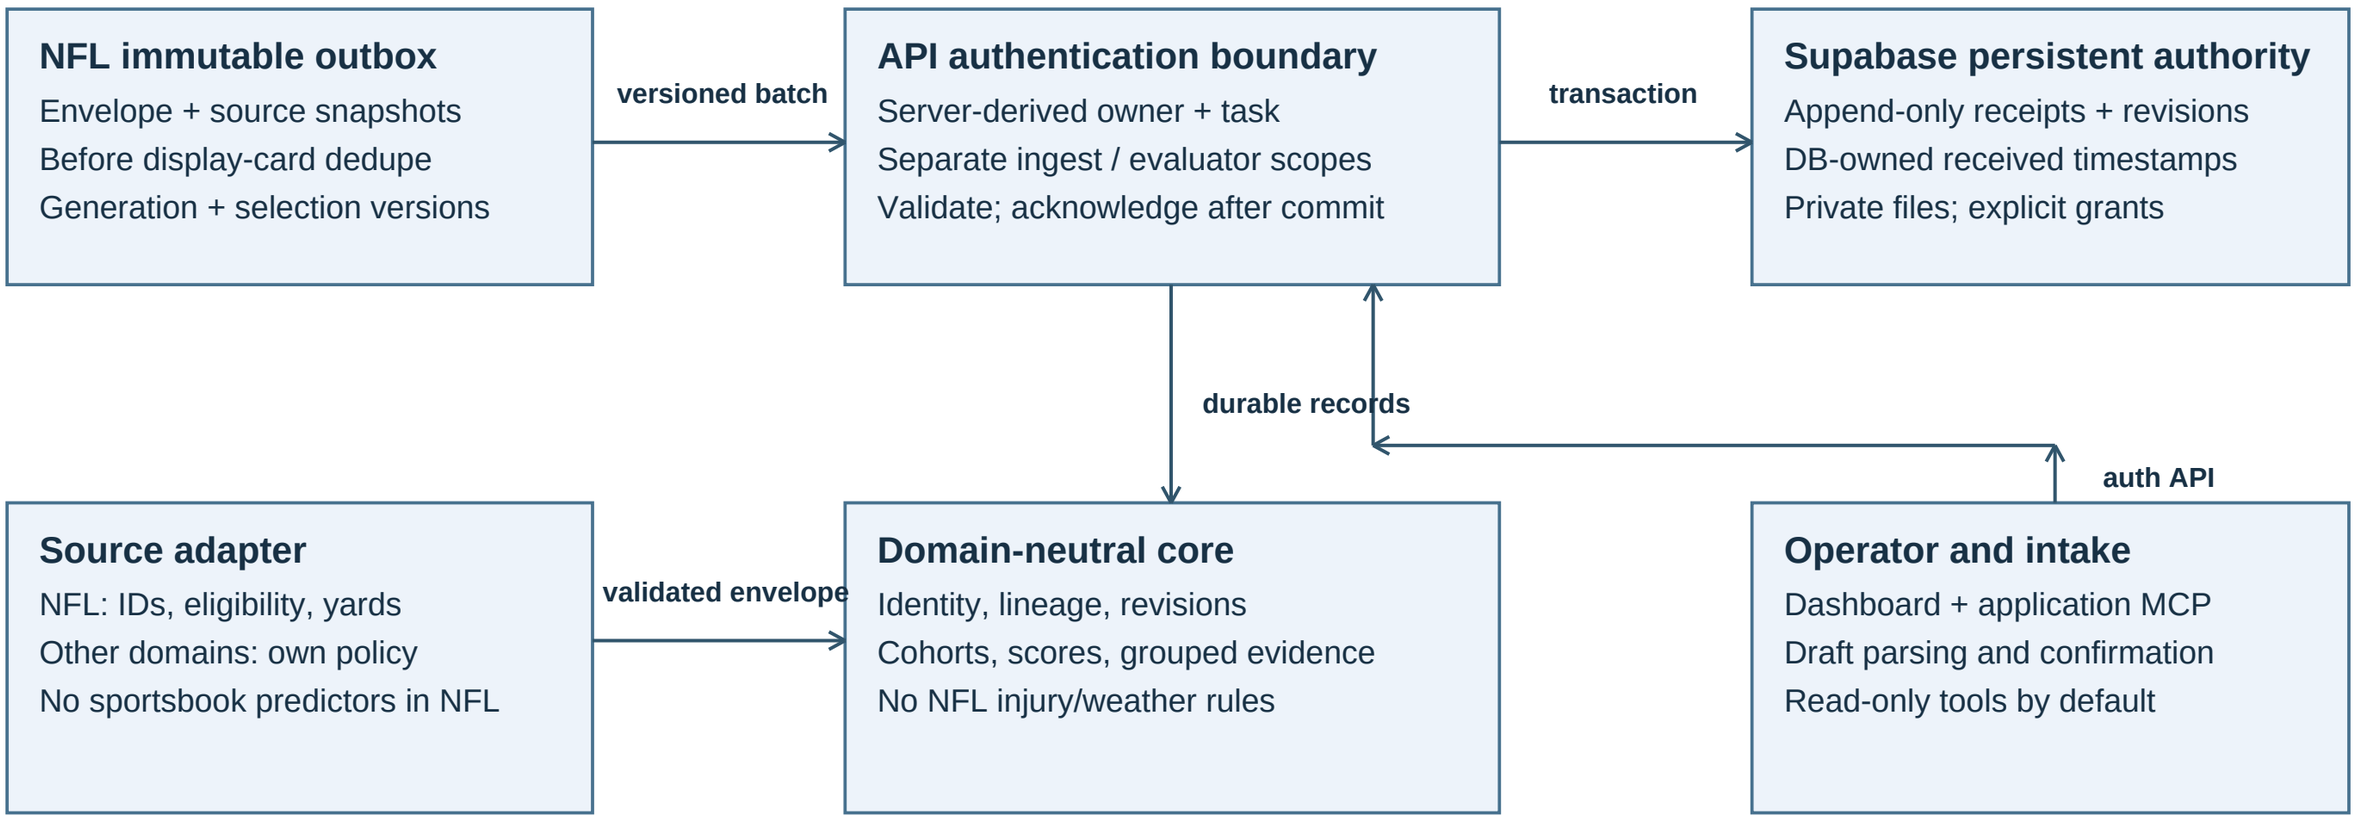

Forecast identity

Contract ID != display card ID != source receipt ID != persisted envelope ID. Preserve each join explicitly.

Original versus correction

New outcome revisions supersede by chain. Old forecasts and evaluation manifests remain unchanged.

Production state

Vercel filesystem and local SQLite are not the hosted store. Supabase schema/API extensions still need build and approval.

Forecast, publication and outcome pipelines

Detailed control flow: freeze evidence before results; distinguish internal capture from observed public availability.

Forecast and publication

1. Fetch as-of sources
2. Generate base -> serve once
3. Journal before dedupe
4. Gate exact generation
5. Publish + verify hashes
6. Export + durable ACK

Failure path before publication

- Unknown status -> quarantine its identifiable row.
- Stale source or hash mismatch -> reject with reason.
- Keep R102 checks and R106 archive protection.

Failure path after publication

- Engine down -> preserve outbox; retry same IDs.
- Conflicting payload -> reject; never overwrite.
- ACK only after commit; alert before kickoff.

Outcome and learning

1. Authoritative FINAL
2. Append revision chain
3. Pin exact cohort/revisions
4. Chronological evaluation
5. Shadow candidate
6. NFL gate + owner review

Three clocks: source availability -> forecast creation -> internal/database observation. A later import cannot prove earlier capture. Public availability needs a separate deployment/content-hash observation.

Sampling boundary

Pool rungs, MY offers and Game/Week cards are distinct populations. Compare the same declared cohort; copies are not independent samples.

Scheduling and recovery

Actions completion + manual dispatch; cron is best effort. Durable leases/checkpoints recover repeated or killed jobs.

Learning and adoption states

Statistical and serving authority stay separate. Uploading evidence or fitting a model is not permission to use it.

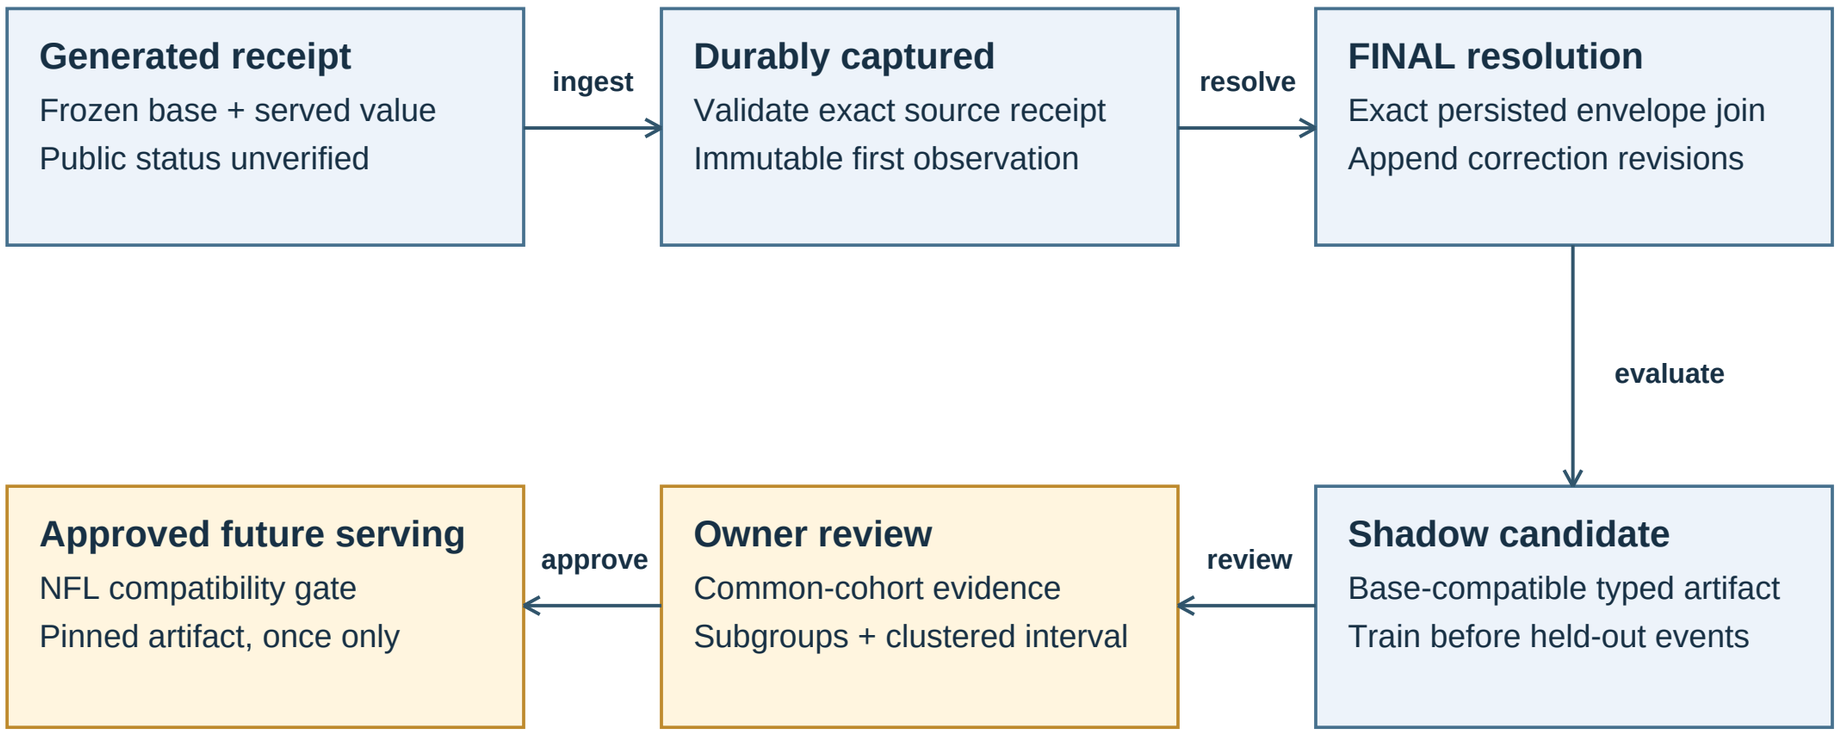

Equation contract

Generic Platt: sigmoid(a x confidence + b). NFL: sigmoid(a + b x logit(immutable base)), then clamp and round. These are not interchangeable.

Timing and rollback

Current MY policy needs >=100 legs and >=2 held-out week folds: normally >=3 receipt weeks. R127 separately needs >=2 new graded weeks. Rollback changes future serving, not history.

> Do not refit to Week 5 or one TNF result. Candidate population, base model, artifact hash, evaluation revisions and effective version must all remain traceable.

Operator and chat journeys

User action -> implementation -> visible result -> observable acceptance. Mobile dashboard is the fallback if chat cannot handle images.

1 / Is learning actually connected?

Open status > Authenticated API + DB check > SHA, durable ingest, job age > UI-01 / OPS-03

Empty, stale and unauthorized states have recovery guidance. MCP connected alone never displays "learning active".

2 / Why this prediction or losing parlay?

Select receipt/card > Canonical trace + cohort > Source, versions, exposure > EVAL-01 / PORTFOLIO-02

Show generated/public, pending/FINAL, original/reduced settlement and shared-loss blast radius without changing the original bet.

3 / Analyze my screenshot

Private upload > Draft parse + confirm IDs > Joined ticket + uncertainty > INTAKE-01/02/03

Ambiguous game/player/line stays unresolved. Parsed text is untrusted; it cannot create verified publication or a FINAL label.

4 / Should we use this improvement?

Compare candidate > Pinned held-out evaluation > Hold reason or approval review > PROMOTE-01 / RETURN-02

Read-only MCP tools inspect evidence. Approval is a separate authorized action; active version means consumed by NFL, not latest proposed.

Questions are not automatically training data. Objective predictions need a defined horizon and authoritative outcome. Subjective approval stays review feedback, not truth.

Parallel team and implementation roadmap

Four execution slots in this session. Speed comes from disjoint ownership and frozen contracts, not skipping evidence or review.

Coordinator

Architecture, contract ownership, integration manifest, one owner question when needed.

NFL lane

Eligibility, generation journal, publisher gate, exporter and approved consumer.

Engine lane

Adapters, authenticated API, durable store, revisions and migration proposal.

Evaluation / UX lane

Cohorts, replay, metrics and operator views; independent QA rotates into a freed slot.

M0 / Foundation	Reproducible source, schema and scoped credentials	1-2 engineer-days
M1 / Prospective capture	Journal before dedupe, exact generation gate	2-3 days
M2 / Durable shadow loop	API, outbox/ACK, persistence, FINAL revisions	3-4 days
M3 / Evaluation and return	Parity, typed artifact, independent NFL gates	3-5 days
M4 / Operator and NFL work	R130, portfolio, R127/weather shadow, MCP/status	2-4 days; overlaps
M5 / M6 / Adopt then extend	Evidence + owner approval; one domain pilot at a time	Observation time extra

Supported MY shadow-to-controlled-return build: roughly 10-18 engineer-days with overlap. Broad screenshot understanding and real financial/F1/soccer/NCAA/election adapters are additional scope. No same-day accuracy promise.

Claude Code is the proposed release coordinator per NFL workflow, not a live connected agent here. Exact SHA handoffs and independent tests connect the sessions. Owner authorizes production migrations, deploys and activation.

Acceptance gates and the owner’s next actions

The architecture becomes real only when the complete round trip is demonstrated on exact, reproducible revisions.

Required end-to-end evidence

- Pending receipt captured before kickoff, independently of display IDs.
- Commit ACK -> exact retry no-op -> conflict rejected.
- FINAL correction linked to the exact persisted envelope.
- NFL/engine identical cohort and served-score parity.
- Earlier-base fit -> shadow candidate -> incompatible return rejected.
- Approved future artifact applied once -> second cycle -> outage/rollback.

Your remaining actions

- No new operator-token setup; keep secret values out of chat.
- If convenient, confirm SUPABASE_URL and a supported backend secret key are configured; they may already exist.
- Approve reviewed migration and release only after test evidence.
- Later choose portfolio tradeoff and any candidate activation.
- No new paid provider or automatic promotion is needed to start.

Current verified foundation

Supabase project + 4 RLS-enabled tables; migration history and advisor results. Operator token present per owner.

Current limits

Hosted end-to-end auth/DB round trip unverified. Helper grants need review; migration IDs differ. Remote checks are restricted here.

Local implementation evidence

Producer-parity and journal tests passed locally. This is not hosted persistence, production ingestion, candidate adoption or proven accuracy gain.

Source package: TECHNICAL_PLAN.md / BACKLOG.md / CSV / backlog-source.mjs / render.mjs / editable architecture.html. 12 epics / 31 stories / 93 tasks / 92 acceptance criteria.

Primary references: vercel.com/docs/frameworks/backend/fastapi / supabase.com/docs/guides/api/securing-your-api / GitHub scheduled-workflow docs. Technical plan links exact source baselines and evidence.